Tianjian Li

Baltimore, MD | tli104@jhu.edu | Personal Website | Github | Google Scholar

Research Interests

My research lies at the intersection of Natural Language Processing and Machine Learning. In particular, I focus on **data engineering** for language models (LMs), including filtering low-quality data, balancing data across diverse sources, and enhancing the diversity of training data for LMs.

Education

Johns Hopkins University, Ph.D. in Computer Science — Aug 2024 – Present

- Advisor: Daniel Khashabi
- Research Focus: Natural Language Processing

Johns Hopkins University, M.S. in Computer Science — Aug 2022 - May 2024

- Advisor(s): Kenton Murray, Philipp Koehn, Daniel Khashabi
- Research Focus: Natural Language Processing

New York University, B.A. Joint Major in Mathematics and Computer Science — Aug 2017 - May 2021

Industry Experience

Research Scientist Intern, Meta Fundamental AI Research (FAIR) – Bellevue, WA — May 2025 – Dec 2025

- Advisor(s): Tianlu Wang, Jack Lanchantin, Jason Weston

Machine Learning Engineer, Baidu Inc. – Beijing, China — Aug 2021 – Jan 2022

First-Authored Publications

Jointly Reinforcing Diversity and Quality in Language Model Generations [Link] [Code]
Tianjian Li, Yiming Zhang, Ping Yu, Swarnadeep Saha, Daniel Khashabi, Jason Weston, Jack Lanchantin, Tianlu Wang.
Preprint 2025.

SimpleMix: Frustratingly Simple Mixing of Off- and On-policy Data in Language Model Preference Learning [Link]
Tianjian Li, Daniel Khashabi.
ICML 2025.

Upsample or Upweight: Balanced Training on Heavily Unbalanced Datasets [Link][Video]
Tianjian Li, Weiting Tan, Haoran Xu, Kenton Murray, Daniel Khashabi.
NAACL 2025.

Error Norm Truncation: Robust Training in the Presece of Data Noise in Text Generation Models [Link][Code][Video]
Tianjian Li, Haoran Xu, Philipp Koehn, Daniel Khashabi, Kenton Murray.
ICLR 2024. Spotlight Presentation ($367/7304 \approx 5\%$).

Why Does Zero-shot Cross-lingual Transfer Fail? An Explanation and a Solution [Link]
Tianjian Li, Kenton Murray.
ACL 2023 Findings.

Non-First Authored Publications

The Flaw of Averages: Quantifying Uniformity of Performance on Benchmarks [Link]
Arda Uzunoglu, **Tianjian Li**, Daniel Khashabi. Preprint 2025.

The Translation Barrier Hypothesis: Multilingual Generation with Large Language Models Suffers from Implicit Translation Failure [Link]

Niyati Bafna, **Tianjian Li**, Kenton Murray, David R. Mortensen, David Yarowsky, Hale Sirin, Daniel Khashabi.
Preprint 2025.

Benchmarking Language Model Creativity: A Case Study on Code Generation [Link][Code][Video]

Yining Lu, Dixuan Wang, **Tianjian Li**, Dongwei Jiang, Sanjeev Khudanpur, Meng Jiang, Daniel Khashabi.
NAACL 2025.

Verifiable by Design: Aligning Language Models to Quote from Pre-Training Data [Link][Code][Video]

Jingyu Zhang, Marc Marone, Tianjian Li, Benjamin Van Durme, Daniel Khashabi.
NAACL 2025. Oral Presentation.

Technologies

Languages: Python, C++, Java, Shell Script.

Frameworks and Tools: Linux, PyTorch, HuggingFace, Inference Engines (vLLM, SgLang), RL frameworks (verl, slime), Distributed Training (FSDP, DeepSpeed), vim.

Awards

National Olympiad in Informatics Provinces (NOIP) First Prize. 2015

Service

Reviewer for —

- NLP Conferences: ACL (2023, 2024, 2025), EMNLP (2023, 2024, 2025), NAACL (2025), EACL (2024, 2025), COLM (2024, 2025).
- ML Conferences: ICLR (2025, 2026), NeurIPS (2025)

Organizer: Mid-Atlantic Student Colloquium on Speech, Language and Learning (2024)

Teaching

Head Teaching Assistant, JHU CS 601.471/671 NLP: Self-supervised Models Feb 2025 – May 2025

- Taught Lectures on Distributed Training: [Video 1][Video 2]